

Software Compiler Assignment on:

Vector Operations

The Software Compiler course exam is composed by two parts. One is a written test, the other is an homework, to be terminated before course last class. The written test contributes with the 40% to the whole grade, while the homework contributes with the remaining 60%. To pass the whole exam, you must get a pass grade from both the test and the homework.

During the lab classes, we show you the ACSE compiler and the associated assembler [2]. They must be used as starting point for your homework. During the last class, you must present your work, showing your edits, giving a brief demo, and responding at some questions given by course lecturers.

Sources of ACSE can be found on the course site [3]. It is a tarball of the ACSE mercurial [1] repository. You are required to version your code with mercurial. A copy of your edits must be submitted to course lectures in the form of a patch with respect to the version of ACSE you have used as starting point.

Assignment

You are required to modify both the front-end of the compiler and the back-end.

Front-end

Vector operations allows to perform simple operations – e.g. adding, subtracting, ... – to all elements stored in an ACSE array. Figure 1 shows some examples.

You are required to add support to the ACSE front-end for handling the following vector operations:

Addition The `vec_add(c, a, b)` statement perform an element wise addition between arrays `a` and `b`, storing the result into the array `c`, that is:

$$\forall i \mid 0 \leq i < \text{length}(c), c[i] \leftarrow a[i] + b[i]$$

The three arrays must have the same length. If this does not hold, a compile-time error must be raised. Code generator must target the standard ACSE assembly.

<pre>int a[10], b[10], c[10], d[10]; vec_add(c, a, b); vec_sub(d, a, b);</pre>	<pre>int a[10], b[10], c[10]; int cond; vec_write(a); vec_blend(c, a, b, cond);</pre>
--	--

Figure 1: Example of vector operations

Subtraction The `vec_sub(c, a, b)` statement perform an element wise subtraction between arrays `a` and `b`, storing the result into the array `c`, that is:

$$\forall i \mid 0 \leq i < \text{length}(c), c[i] \leftarrow a[i] - b[i]$$

The three arrays must have the same length. If this does not hold, a compile-time error must be raised. Code generator must target the standard ACSE assembly.

Printing The `vec_print(a)` statement print the whole array. Code generator must target the standard ACSE assembly.

Blending The `vec_blend(c, a, b, cond)` allows setting `c` according to expression `cond`. If it holds, then vector `a` is copied into vector `c`. If it doesn't hold, then vector `b` is copied into vector `c`:

$$c \leftarrow \begin{cases} a & \text{if } cond \neq 0 \\ b & \text{if } cond = 0 \end{cases}$$

The three arrays must have the same length, otherwise a compile-time error must be raised. Code generator must target standard ACSE assembly.

Back-end

Peephole optimizations are optimizations applied later in the compilation process. Their goal is to identify groups of subsequent instructions that can be substituted with another, more efficient, instruction.

Let `a`, `b`, `c`, three labels in the data segment associated to `SPACE` directives. They can be seen as a representation of ACSE arrays in assembler language. Suppose that the following instructions are added to the ACSE assembly language:

- **VADD c a b**: interprets memory locations identified by `a`, `b`, and `c` as three 4 elements arrays, and performs element-wise addition of `a` and `b`, storing the result in `c`
- **VSUB c a b**: interprets memory locations identified by `a`, `b`, and `c` as three 4 elements arrays, and performs element-wise subtraction of `a` and `b`, storing the result in `c`

...	
MOVA R1 L1	
ADD R2 R0 (R1)	
MOVA R1 L2	
ADD R3 R0 (R1)	
ADD R1 R2 R3	
MOVA R2 L0	...
ADD (R2) R0 R1	VADD L0 L1 L2
...	...
(a) Source	(b) Optimized

Figure 2: Folding an addition

You are required to implement a simple peephole optimization, looking for instruction sequences in the ACSE assembler that can be implemented using `VADD` or `VSUB`. Figure 2 shows an example.

If the pattern reported in Figure 2(a) is repeated 4 times without interruptions, then the optimizer can fold the addition and generate the code reported in Figure 2(b).

Your optimizer must be a new ACSE tool, like `asm` – which can be used as starting point. The tool must read an input text assembly, apply the optimization where possible, and write in a new file the output assembly.

Creating and Applying Patches

The ACSE distribution is already provided as a mercurial repository. See [4] for the basic mercurial commands.

Assuming that your work has been committed with revision `X`, to generate a patch use the following command:

```
hg diff -r 0 -r X
```

Assuming the patch has been saved in the file `patch.diff`, use the following command to apply it to a freshly unpacked ACSE source tree:

```
hg import patch.diff -m "Applied patch"
```

The submitted project must successfully pass the above process, i.e. before submitting your patch check that it can be applied to a freshly unpacked ACSE source tree including compilation and testing process where applicable.

References

[1] Mercurial. <http://mercurial.selenic.com>, 2011.

- [2] A. Di Biagio and G. Agosta. Advanced Compiler System for Education. <http://compilergroup.elet.polimi.it>, 2008.
- [3] Formal Languages and Compilers Group. Software Compilers. <http://compilergroup.elet.polimi.it>, 2010.
- [4] J. Spolsky. Hg Init: a Mercurial Tutorial. <http://hginit.com>, 2011.